

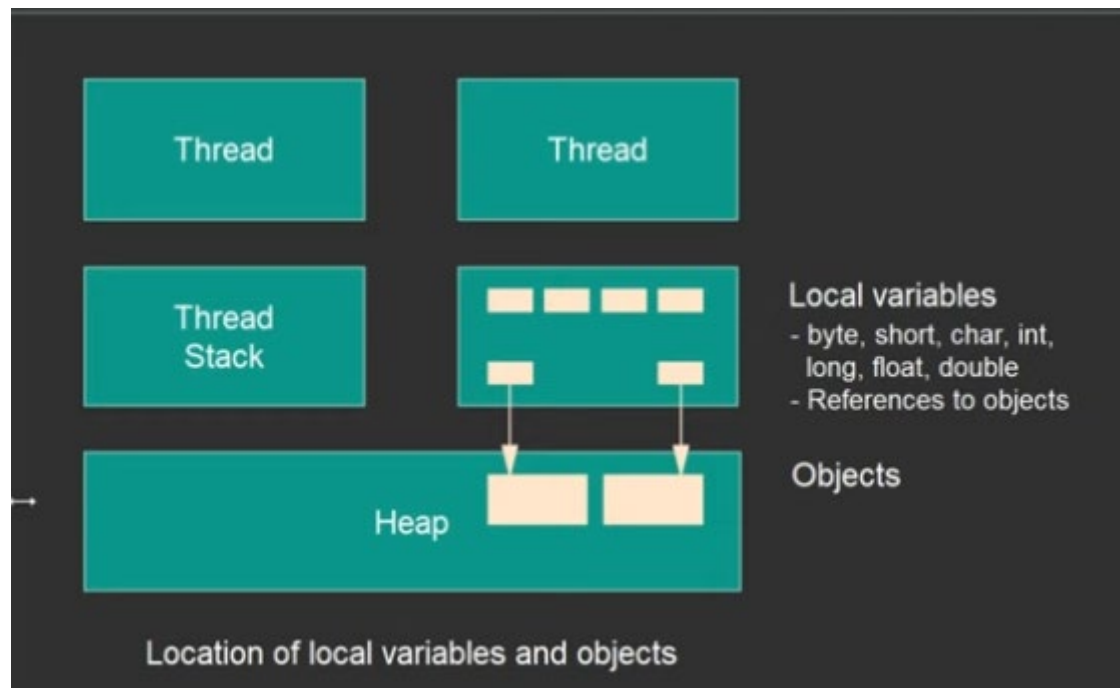


*Presents*

# **Concurrency II**

# Java Thread Memory

- ▶ Java threads have their own stacks but share heap memory
  - ✓ For objects on the heap, each thread has its own reference to that object on its stack



# Thread Objects

- ▶ Objects created by a runnable object are only accessible to that object
  - ✓ The object is created on the shared heap
  - ✓ But only the thread's stack has a reference to it
  - ✓ No other threads can access the object because they don't have a reference to it.
- ▶ Consider the following test object

```
1
2
3 public class MyObject {
4     String name = null;
5
6     public MyObject(String s) {
7         this.name = s;
8     }
9
10 }
```

# Thread Objects

- ▶ Each Runnable object creates its own copy

```
public class Task implements Runnable {  
  
    String name = null;  
    MyObject obj = new MyObject(this.name);  
  
    public Task(String n) {  
        this.name = n;  
    }  
  
    @Override  
    public void run() {  
        System.out.println("Thread " + this.name + " using object " + obj);  
    }  
  
}
```

# Thread Objects

- ▶ Running two threads shows the addresses of the MyObject objects are different

```
public class Runner {  
  
    public static void main(String[] args) {  
        Thread t1 = new Thread(new Task("Alpha"));  
        Thread t2 = new Thread(new Task("Beta"));  
        t1.start();  
        t2.start();  
    }  
}
```

Console X

Run> Runner (8) [Java Application] C:\tools\java\Java18\bin\javaw.exe (Jul. 9, 2022, 11:17:24 a.m. – 11:17:24 a.m.)

Thread Alpha using MyObject Unshared.MyObject@b8bf527

Thread Beta using MyObject Unshared.MyObject@3d00d211

# Shared Object

- ▶ If the object exists prior to the creation of the runnable object
  - ✓ Then the runnable object can be passed a reference to the object
  - ✓ Then all references in the threads resolve to the same object
  - ✓ Even though each thread has its own copy of the reference, they still share the object their references point to.

# Shared Object

```
public class SharedTask implements Runnable{

    String name = null;
    MyObject myObj = null;

    public SharedTask(String n, MyObject obj) {
        this.name = n;
        this.myObj = obj;
    }

    @Override
    public void run() {
        System.out.println("Thread " + this.name + " using object " + this.myObj);
    }
}
```

# Shared Objects

- ▶ Each runnable object gets a reference to an existing object

```
3 public class SharedRunner {  
4  
5     public static void main(String[] args) {  
6         MyObject singleton = new MyObject("shared object");  
7  
8         Thread t1 = new Thread(new SharedTask("one", singleton));  
9         Thread t2 = new Thread(new SharedTask("two", singleton));  
10        t1.start();  
11        t2.start();  
12    }
```

Problems @ Javadoc Declaration Console X

<terminated> SharedRunner [Java Application] C:\tools\java\jdk-17.0.2\bin\javaw.exe (Feb. 27, 2022, 2:01:16 p.m. – 2:01:16 p.m.)

Thread two using object threads.MyObject@e0c18a6 shared object

Thread one using object threads.MyObject@e0c18a6 shared object



# Race Conditions

- ▶ A race condition occurs when more than one thread tries to access a shared resource
  - ✓ Not a problem if all the threads are just reading the same resource
  - ✓ Modifying or writing data in a shared resource may result in corruption
- ▶ Result of how the threads are interleaved
  - ✓ Some instructions from one thread may overlap with instructions from another thread
- ▶ The problem is that because race conditions are environment dependent
  - ✓ They are hard to test for
  - ✓ They may be transient (not occur every time)

# Race Example

- ▶ Suppose we have a shared counter and two threads are trying to update it.

```
int counter = 0;  
int newVal = 0
```

<code>newVal1 = counter + 1;</code>	<code>// Thread one reads the counter and increments</code>
<code>newVal2 = counter + 1;</code>	<code>// Thread two reads and increments</code>
<code>counter = newVal1</code>	<code>// Thread 1 updates counter</code>
<code>counter = newVal2</code>	<code>// Thread 2 overwrites the value just saved</code>

# Race Types

## ► Check-Then-Act race condition

- ✓ Logic executed by a thread depends on the value of a shared variable
- ✓ A thread may read a stale value and perform the wrong action

## ► Read-Modify-Write

- ✓ This is the example we just saw
- ✓ Multiple threads modifying the same resource
- ✓ Because of the interleaving of reads and writes, the value becomes incorrect

# Avoiding Race Conditions

- ▶ Restrict access to critical sections of code where race conditions could occur
- ▶ Java uses the “synchronized” keyword
  - ✓ Ensures that the resource identified uses atomic operations
  - ✓ This means essentially only one thread at a time can access the resource
- ▶ When a thread enters a synchronized block all other threads trying to execute that block must wait
  - ✓ Once the thread executes the block, the next thread can have access to the block
- ▶ This is illustrated in the lab

# Semaphores

- ▶ The problem with using the synchronized keyword is that threads block waiting for their turn
  - ✓ This can create slowdowns in processing
  - ✓ The threads are idle while waiting to use the synchronized resource
- ▶ A semaphore is an integer variable shared by the threads
  - ✓ Semaphores regulate access to a shared resource
  - ✓ Eliminates busy waiting
  - ✓ Can be binary taking the values 0 and 1
  - ✓ Or can be a count of available resources

# Semaphores

- ▶ Two operations
- ▶ Wait or wait for:
  - ✓ Take a resource and decrements the semaphore
  - ✓ If the semaphore is 0, then the thread sleeps until one is available
- ▶ Signal:
  - ✓ Releases a resource and increments the semaphore
  - ✓ If there are sleeping threads, a waiting thread is woken up

# Java Semaphore

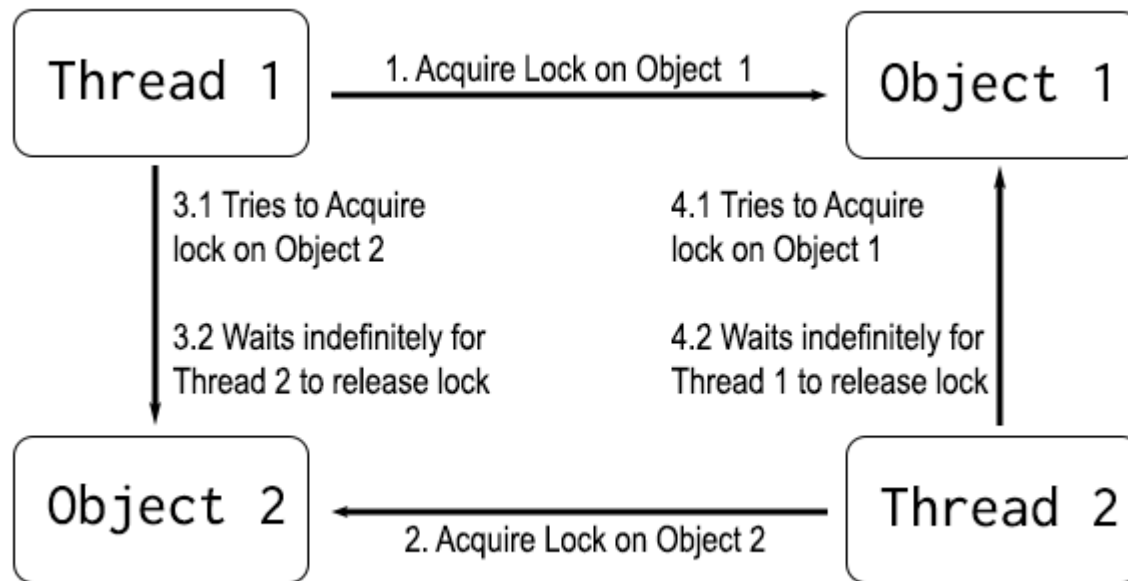
- ▶ A counting semaphore.
  - ✓ Maintains a set of permits
  - ✓ Each `acquire()` blocks if necessary, until a permit is available, and then takes it
  - ✓ Each `release()` adds a permit, potentially releasing a blocking acquirer.
  - ✓ However, no actual permit objects are used; the Semaphore just keeps a count of the number available and acts accordingly.

# Mutex

- ▶ Similar to a binary semaphore
  - ✓ However, it is more like a lock in that a thread can take ownership of a mutex
  - ✓ It then has exclusive use of the resource until it releases the mutex
- ▶ The synchronized keyword in Java essentially is a mutex
- ▶ We can emulate a mutex in Java using a semaphore



# Mutex Deadlock



# Atomic Access

- ▶ An atomic action happens all at once.
  - ✓ It cannot stop in the middle
  - ✓ it either happens completely, or it doesn't happen at all
  - ✓ No side effects of an atomic action are visible until the action is complete
- ▶ In Java
  - ✓ All primitive types use atomic reads and writes except for longs and doubles
  - ✓ All variables marked volatile also use atomic reads and writes

# Atomic Variables

- ▶ Java provides atomic classes that support atomic operations on single variables
  - ✓ They have `get()` and `set()` methods that operate like atomic reads and writes

# Atomic Variables Example

- ▶ Consider the following Counter class

```
class Counter {  
    private int value = 0;  
  
    public void increment() {this.value++;}  
    public void decrement() {this.value--;}  
    public int value() { return this.value;}  
}
```

# Atomic Variables Example

- ▶ We can use a mutex to control access

```
class SynchCounter {  
    private int value = 0;  
  
    public synchronized void increment() {this.value++;}  
    public synchronized void decrement() {this.value--;}  
    public synchronized int value() { return this.value;}  
}
```

# Atomic Variables Example

- ▶ We can get the same effect using an Atomic variable

```
class AtomicCounter {  
    private AtomicInteger value = new AtomicInteger(0);  
  
    public void increment() {this.value.incrementAndGet();}  
  
    public void decrement() {this.value.decrementAndGet();}  
  
    public int value() {return this.value.get();}  
  
}
```

# Concurrent Collections

## ► Memory Consistency Errors

- ✓ Occur when different threads have inconsistent views of what should be the same data
- ✓ Managed by a “happens-before” relationship
- ✓ Guarantee that memory writes by one specific statement are visible to another specific statement

## ► If multiple threads are accessing a collection

- ✓ We may have two threads trying to add a “last” object into the same position while another thread is reading the last object
- ✓ It's not clear what the “last object” will actually be to each thread

# Concurrent Collections

- ▶ **BlockingQueue**
  - ✓ Defines a first-in-first-out data structure that blocks or times out on attempts to add to a full queue, or retrieve from an empty queue.
- ▶ **ConcurrentMap**
  - ✓ Defines atomic operations for maps.
- ▶ There are other more specialized collections
- ▶ Concurrency issues are managed by the collection
  - ✓ Recall that the Executor Service uses a BlockingQueue



# The Thread Problem

- ▶ Threads run asynchronously
  - ✓ The spawning thread does not wait for the child thread to finish
  - ✓ The Runnable run() method returns a void
  - ✓ Where would a return value be returned to?
- ▶ Callbacks
  - ✓ One approach is to include a function that is called when the thread is done
  - ✓ Named a “call back” (“Call me when you finish”)
  - ✓ Standard pattern in asynchronous programming

# Callback Issues

- ▶ Error handling
  - ✓ No clean mechanism for handling errors in a callback
  - ✓ Also, how does a callback handle an error in the thread?
- ▶ Nested callbacks
  - ✓ If we have a chain of threads, each of which has a callback
  - ✓ Then we have a nested chain of callbacks
  - ✓ Can be very difficult to manage at runtime
  - ✓ Very hard to debug
  - ✓ What thread manages the return values from the callback,
- ▶ Blackboard pattern
  - ✓ Callbacks all update a common data structure

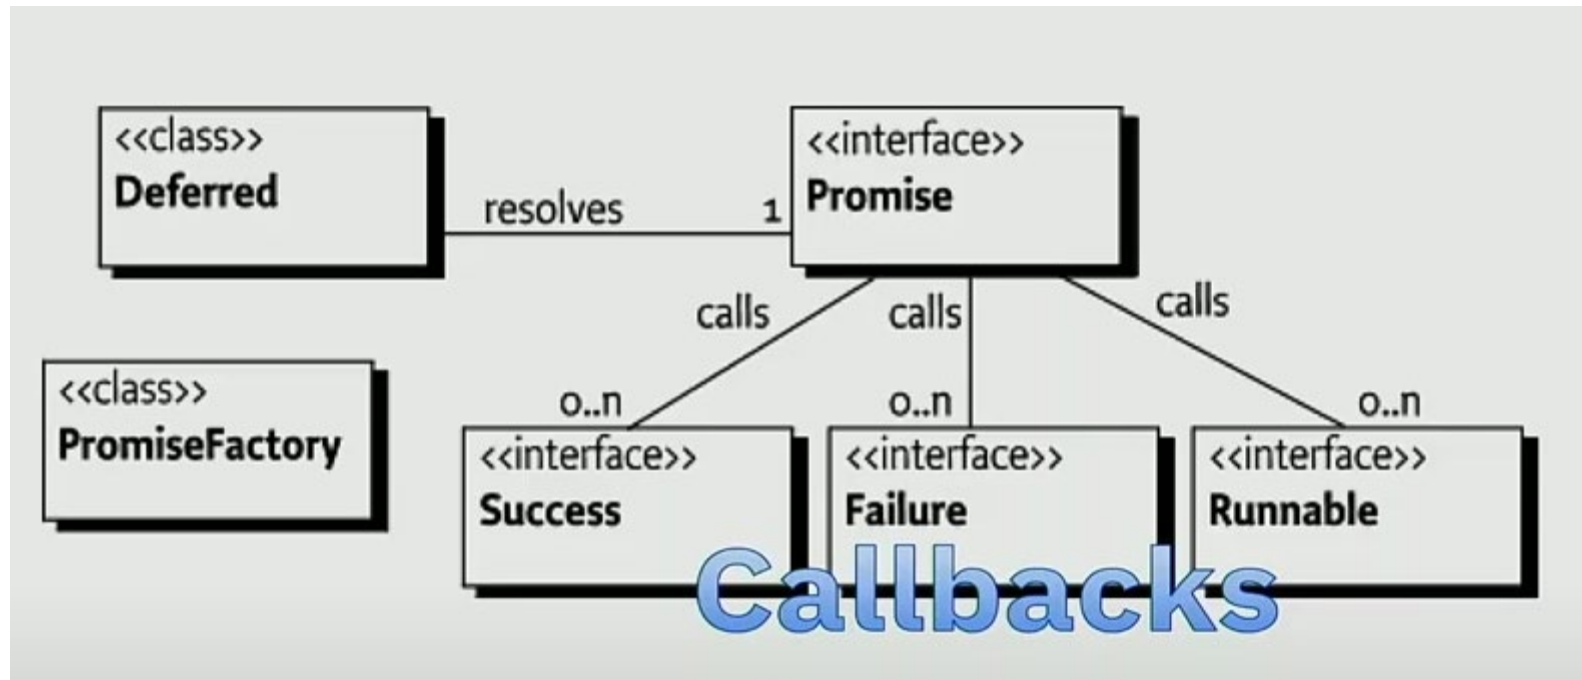
# Callbacks in Java

- ▶ Callbacks are handled in various languages by passing a function pointer to the callback function
  - ✓ Java doesn't have function pointers
- ▶ Java emulates this by using a callback interface
  - ✓ The callback object implements the callback interface
  - ✓ When the thread starts, it is passed a callback object
  - ✓ When an event that requires a callback happens, the callback method is invoked on the callback object

# Promises

- ▶ A Promise represents a proxy
- ▶ Stands in for for a value not known when the promise is created
- ▶ Promises are associated to an asynchronous action's eventual success value or failure
- ▶ A promise is a placeholder for the result of an ongoing asynchronous operation
- ▶ A promise can be resolved or unresolved

# Promises



# Promises

- ▶ The task creates a deferred object that will be resolved to a promise
  - ✓ A promise to be filled in after the task is completed
- ▶ After the task completes the promise is resolved to one of the outcomes in the interface
  - ✓ The outcomes are callbacks, just like we saw in the previous section

# Questions

